



# The Zwicky Transient Facility Alert Distribution System

Maria T. Patterson<sup>1</sup> , Eric C. Bellm<sup>1</sup> , Ben Rusholme<sup>2</sup> , Frank J. Masci<sup>2</sup> , Mario Juric<sup>1</sup> , K. Simon Krughoff<sup>3</sup> ,  
V. Zach Golkhou<sup>1,4,6</sup> , Matthew J. Graham<sup>5</sup> , Shrinivas R. Kulkarni<sup>5</sup> , and George Helou<sup>2</sup>

Zwicky Transient Facility Collaboration

<sup>1</sup> DIRAC Institute, Department of Astronomy, University of Washington, 3910 15th Avenue NE, Seattle, WA 98195, USA; [mtpatter@uw.edu](mailto:mtpatter@uw.edu)

<sup>2</sup> IPAC, California Institute of Technology, 1200 E. California Boulevard, Pasadena, CA 91125, USA

<sup>3</sup> LSST Project Office, 950 N. Cherry Avenue, Tucson, AZ 85719, USA

<sup>4</sup> The eScience Institute, University of Washington, Seattle, WA 98195, USA

<sup>5</sup> Division of Physics, Mathematics, and Astronomy, California Institute of Technology, Pasadena, CA 91125, USA

Received 2018 September 14; accepted 2018 October 10; published 2018 November 27

## Abstract

The Zwicky Transient Facility (ZTF) survey generates real-time alerts for optical transients, variables, and moving objects discovered in its wide-field survey. We describe the ZTF alert stream distribution and processing (filtering) system. The system uses existing open-source technologies developed in industry: Kafka, a real-time streaming platform, and Avro, a binary serialization format. The technologies used in this system provide a number of advantages for the ZTF use case, including (1) built-in replication, scalability, and stream rewind for the distribution mechanism; (2) structured messages with strictly enforced schemas and dynamic typing for fast parsing; and (3) a Python-based stream processing interface that is similar to batch for a familiar and user-friendly plug-in filter system, all in a modular, primarily containerized system. The production deployment has successfully supported streaming up to 1.2 million alerts or roughly 70 GB of data per night, with each alert available to a consumer within about 10 s of alert candidate production. Data transfer rates of about 80,000 alerts/minute have been observed. In this paper, we discuss this alert distribution and processing system, the design motivations for the technology choices for the framework, performance in production, and how this system may be generally suitable for other alert stream use cases, including the upcoming Large Synoptic Survey Telescope.

**Key words:** astronomical databases: miscellaneous – instrumentation: miscellaneous – surveys

**Online material:** color figures

## 1. Introduction

The Zwicky Transient Facility (ZTF; Bellm et al. 2018; Graham et al. 2018) distributes a live stream of its detections of transient astronomical events at a rate of about 1 million alerts per night. Contemporaneous observations by different types of telescopes and instruments at a range of wavelengths is essential to understand these astrophysical phenomena (e.g., Gal-Yam et al. 2014; Cao et al. 2015; Cenko et al. 2015). In the past, a major limiting factor in securing followup observations of transient events was the delay between the initial detection of an object by one telescope and reporting to other telescopes.

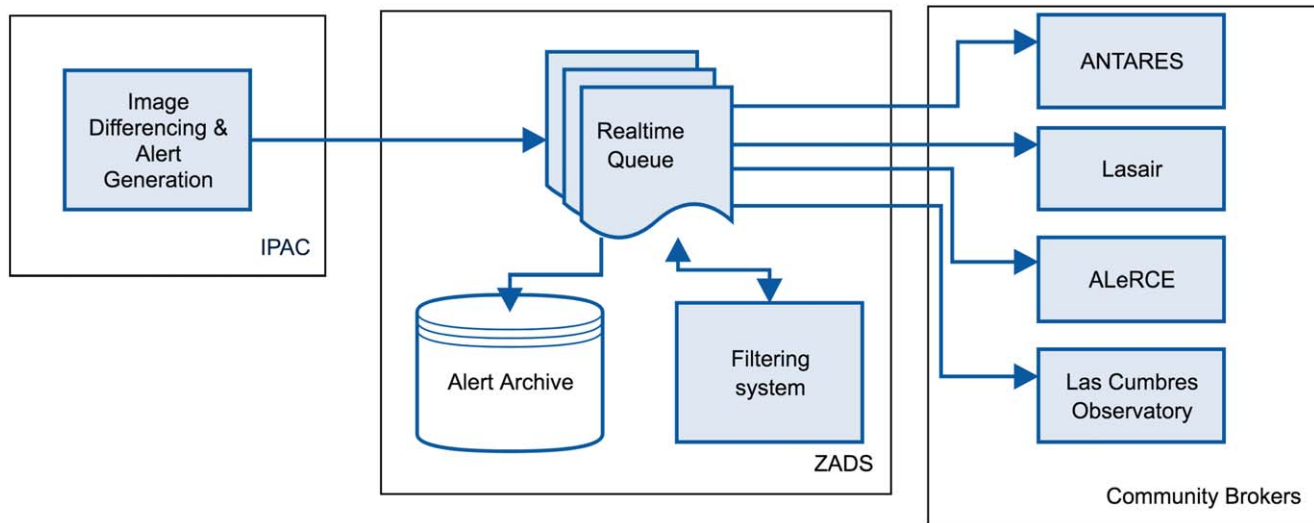
The use of electronic distribution mechanisms such as the GRB Coordinates Network (GCN), the Astronomer’s Telegram (ATel; Rutledge 1998), IAU Circulars, the Transient Name Server,<sup>7</sup> and services relaying Virtual Observatory Event (VOEvent; Williams & Seaman 2006) messages have significantly alleviated the dissemination bottleneck. These mechanisms provide fast dissemination of transient alerts, but the data volume to date has been

manageable enough for manual and visual processing upon receipt. For reference, the 4 Pi Sky collaboration maintains a database of VOEvent alert events going back to 2014 April that includes about 2.2 million alerts as of the writing of this paper (see Staley & Fender 2016a, 2016b). The largest volume of these alerts are reports of *Swift* re-pointings broadcast by the GCN, but they also include photometric alerts from the ASAS-SN and Gaia collaborations among others.

However, in the era of ZTF and upcoming Large Synoptic Survey Telescope (LSST; Ivezić et al. 2008), the bottleneck to discovery will likely shift from distribution to the ability to quickly isolate and prioritize detections of interest amid the flood of alert events: The volume of VOEvents included in the 4 Pi Sky database from all sources is about two nights of alerts from ZTF. Rather than only reporting events likely to be new explosive extragalactic transients, these surveys will stream *all* sources that are above a specified detection threshold in the difference image, whether they are likely due to transients, variable stars, or moving objects. This approach allows the science user the greatest flexibility in identifying candidates of interest for their individual science program. However, these

<sup>6</sup> Moore-Sloan, WRF, and DIRAC Fellow.

<sup>7</sup> <https://wis-tns.weizmann.ac.il/>



**Figure 1.** Schematic overview of ZTF alert flow.

(A color version of this figure is available in the online journal.)

alert streams place greater demands on network bandwidth and require new software systems to filter and aggregate sources of interest. Robust and scalable real-time processing capabilities for automated filtering and analysis of these alerts and subsequent integration capabilities with distribution mechanisms is therefore critical for enabling fast follow-up action and observations of time sensitive astronomical activity.

For ZTF project pipeline needs, we describe here the technologies used by the ZTF Alert Distribution System (ZADS) for streaming these transient alert data, the motivations around the choices for the framework, the performance of the system, and how this system may be scaled for larger astronomical data stream needs.

## 2. ZADS Overview and Goals

ZADS is the near-real-time streaming platform that allows fast access to and filtering of the data products from ZTF’s moving and varying object detection pipeline. ZADS is a component of the ZTF Science Data System (Masci et al. 2018) that obtains alerts from the image differencing object detection and alert generation process and delivers them to downstream brokers and science users (Figure 1). The goal of ZADS is to support the availability of science quality alerts from ZTF within 20 minutes of observation.

Given the high detection alert rate of ZTF, ZADS must be robust enough to support the streaming of a relatively large amount of data—over 1 million alerts per night, each with message size of around 60 KB for a total volume over 70 GB of nightly data. The alerts must be moved out of IPAC’s system and support delivery to the alert archive housed at University of Washington, to a public facing cloud system, and to

downstream community event brokers with no data loss. The system currently feeds four main event broker systems which further distribute and provide access to science users: the Arizona NOAO Temporal Analysis and Response to Events System (ANTARES; Narayan et al. 2018); Lasair,<sup>8</sup> a UK-based broker; an initiative based in Chile called Automatic Learning for the Rapid Classification of Events (ALeRCE)<sup>9</sup>; and Las Cumbres Observatory’s Make Alerts Really Simple (MARS) project.<sup>10</sup> Translation to a VOEvent stream over VTP or automating the distribution of Gamma-ray Burst Coordinates Network (GCN) circulars or ATELS has not been attempted by ZADS, but automatic reporting of candidate supernovae to the Transient Name Server (TNS) is planned.

The ZADS platform is key to enabling real-time astronomical time domain science, including the discovery of young supernovae, the detection of near-Earth asteroids, the identification of stellar variables, and the search for electromagnetic counterparts to gravitational wave sources.

## 3. Technologies for ZADS

### 3.1. Alert Formatting

A new bottleneck to discovery in the flood of alert events is the ability to quickly prioritize interesting alerts, necessitating real-time filtering capabilities for the automated detection of objects of interest. The primary alert format used by a number of projects follows the current IVOA standard VOEvent 2.0 (Seaman et al. 2011), which is an XML-based, semi-structured

<sup>8</sup> <http://lasair.roe.ac.uk/>

<sup>9</sup> <http://alerce.science/>

<sup>10</sup> <https://Mars.lco.global/>

**Table 1**  
Partial Schema of ZTF Avro Alert Packets

Field	Type	Contents
objectId	Long	Unique identifier for this object
candid	Long	Unique identifier for the subtraction candidate
candidate	ztf.alert.candidate	Nested schema record including fields prefixed by “candidate” below
candidate.fid	Int	Filter ID (1 = g; 2 = r; 3 = i)
candidate.ra	Double	Right ascension of candidate; J2000 [deg]
candidate.dec	Double	Declination of candidate; J2000 [deg]
candidate.magpsf	Float	Magnitude from PSF-fit photometry [mag]
candidate.distnr	Float or null	Distance to nearest source in reference image PSF-catalog within 30" [pixels]
candidate.magnr	Float or null	Magnitude of nearest source in reference image PSF-catalog within 30" [mag]
candidate.classtar	Float or null	Star/galaxy classification score of candidate from SExtractor
candidate.rb	Float or null	RealBogus quality score; range is 0–1 where closer to 1 is more reliable
prv_candidate	Array of ztf.alert.candidate	Associated alert candidate records for the last 30 days of history
cutoutScience	ztf.alert.cutout or null	Cutout of the science image
cutoutTemplate	ztf.alert.cutout or null	Cutout of the co-added reference image
cutoutDifference	ztf.alert.cutout or null	Cutout of the resulting difference image

**Note.** ZTF uses nested schemas to organize the data in the alert packet. `ztf.alert` is the top-level namespace, the contents of which are shown above. `ztf.alert` relies on nested schemas `ztf.alert.candidate`, `ztf.alert.prv_candidate`, and `ztf.alert.cutout`. The `prv_candidate` field contains an array of one or more previous subtraction candidates at the position of the alert. These are obtained by a simple cone search at the position of the alert candidate on the last 30 days of history. If there are no previous candidates or upper limits, this field is null. The fields for an individual `prv_candidate` are nearly identical to `candidate` except for the omission of the PS1 matches, previous detection history, and reference image information.

data format that is appropriate for both human interaction and machine-readable scenarios. The XML format typically results in an alert package of larger size compared to other encoding schemes given its redundant tags, and filtering of VOEvents requires significant parsing. Given the needs of the system, the Zwicky Alert Distribution System uses a more structured data format that is smaller and more suitable for fast parsing and filtering.

ZADS transmits alerts using Apache Avro,<sup>11</sup> a binary format developed in the open source Hadoop ecosystem that uses JSON-based schemas in its serialization framework. Using Avro allows for a smaller sized alert message, with on average a factor of six smaller size compared with XML, and also faster serialization and deserialization, with a speedup observed to be approximately 40 times (Maeda 2012). ZTF Avro alerts are approximately half the size of the same alert content in XML format with schema embedded and one-third the size without schema. When packaging many alerts in the same .avro file, only one schema is needed for all alerts. Avro alerts are two-thirds the size of compressed XML.

Avro client libraries are available for many major programming languages, including Python 2 and 3. However, we have found that these libraries can vary in quality, at least for Python 3, particularly with a trade-off between speed and functionality.

The ZTF alert schemas are defined in JSON documents in which the alert contents are described as fields with a name, a data type (e.g., string, int, long, array, null), and an optional

documentation string for human reading. The data types are strictly adhered to and are correctly interpreted when deserialized by the data receiver. The schemas can also define and allow for serializing fields as type “bytes,” which allows for the inclusion of binary format data such as individual files. This makes it possible to embed cutout FITS images of the science, reference, and difference images of detected sources in a much more compact way than can be done with the current VOEvent standard. The schemas used by the writer (ZTF project) and the reader (downstream brokers) can also be different, allowing for the easy addition or removal of fields upon receipt.

ZTF makes use of the ability to “nest” schemas in building the alert format, which offers some advantages over the current data model of VOEvent. The main contents of the alert packets fall under the namespace `ztf.alert` as shown in Table 1. Below `ztf.alert` are the primary components including fields for object identification (`ztf.alert.objectId`), the alert candidate detection data (`ztf.alert.candidate`), and image cutouts (each of type `ztf.alert.cutout`). The alert candidate detection history is of type array, (a list of `ztf.alert.prv_candidate` all with the same schema), which can be of arbitrary length and is easily extensible as the detection history grows.

Note that the current ZADS Avro message format and data model was designed to fulfill the needs of the ZTF survey specifically. For example, candidate detection history is formatted in such a way as to allow for the simple processing of lightcurves, and ZTF alerts include postage stamp cutout

<sup>11</sup> <http://avro.apache.org/>

files as opposed to the URI to a location where they can be accessed. By comparison, VOEvent is intended as a more generalized format and attempts to specify a data model suitable for any celestial event use case. Fulfilling the needs of ZTF stream and merging its alert structure with the more broadly applicable VOEvent could provide a roadmap for evolving IVOA standards in preparation for alert streams from large scale surveys such as the LSST.

### 3.2. Alert Message Distribution

For distributing alerts, ZADS utilizes Apache Kafka, a logging system or messaging queue reinvented as a distributed data streaming platform (Kreps et al. 2011). Kafka was originally developed by LinkedIn and later open sourced through the Apache Software Foundation. Kafka is used for large scale, high throughput and low latency data pipelines, has been deployed for industry applications transmitting over 1 trillion messages per day, and can be scaled easily by deploying a cluster of Kafka “brokers.” ZADS deploys Kafka in a small cluster of three brokers, using Docker,<sup>12</sup> a software containerization platform.

Messages are sent to Kafka message queues by “producers” and separated into partitions within a “topic,” which can be thought of as a named single stream which can be subscribed to or a single log that can be sequentially read from. The alert messages are ordered in the partitions and given an offset and a timestamp. Downstream “consumers” subscribe to individual topics. Instead of keeping one copy of a stream per consumer, Kafka instead tracks the latest message offset read from a topic, making it scale easily with the addition of new consumers. This offset tracking method provides the advantage that the offsets can be changed so that consumers can seek to previous alert messages and receive alerts sent with timestamps prior to connecting to the stream. This ensures that those listening to the stream will not miss alert messages and also allows consumers to rewind to past offsets in order to reprocess data. This rewind feature is not available in e.g., services deploying the current transfer system of VOEvents, the VO Event Transport Protocol (VTP; Allan et al. 2017)

Alert consumers may wish to copy an entire topic or topics from the ZADS Kafka system to a locally deployed Kafka cluster from which downstream processing can be done. This consumer is essentially a downstream mirror of topics from the original cluster. Kafka’s “MirrorMaker” feature makes this process of replicating streams straightforward, by deploying a consumer of topics from the source cluster and a producer to the target cluster.

In addition to MirrorMaker, there are a variety of tools in the Kafka ecosystem, which we do not as of yet make use of. Kafka Connect allows users to simply define sources of streams

and sinks to external databases and other storage systems. Kafka Streams is an internal pipeline development library, but it does not currently support Python. KSQL enables data processing with a SQL-like language and supports data of type JSON and Avro. However, use with ZTF’s Avro formatted data would also require Kafka’s schema registry system, which we have not yet explored.

Though we have been able to successfully put Kafka into production, the deployment has not been without some degree of difficulty. As a relatively young technology, it is not commonplace, particularly for science use cases, and expertise in the community is somewhat lacking. Confluent, the company founded by the team who first built Kafka, does, however, provide a Slack team for discussion currently with over 6000 users. We have also found that the documentation for Kafka, though extensive and revised for each distribution, can occasionally be rather terse. Authentication and authorization of users and access control is available but difficult to set up. Additionally, accounting support within Kafka, which would allow, e.g., monitoring of data transfer to clients, is unavailable as an open source product.

### 3.3. Filtering

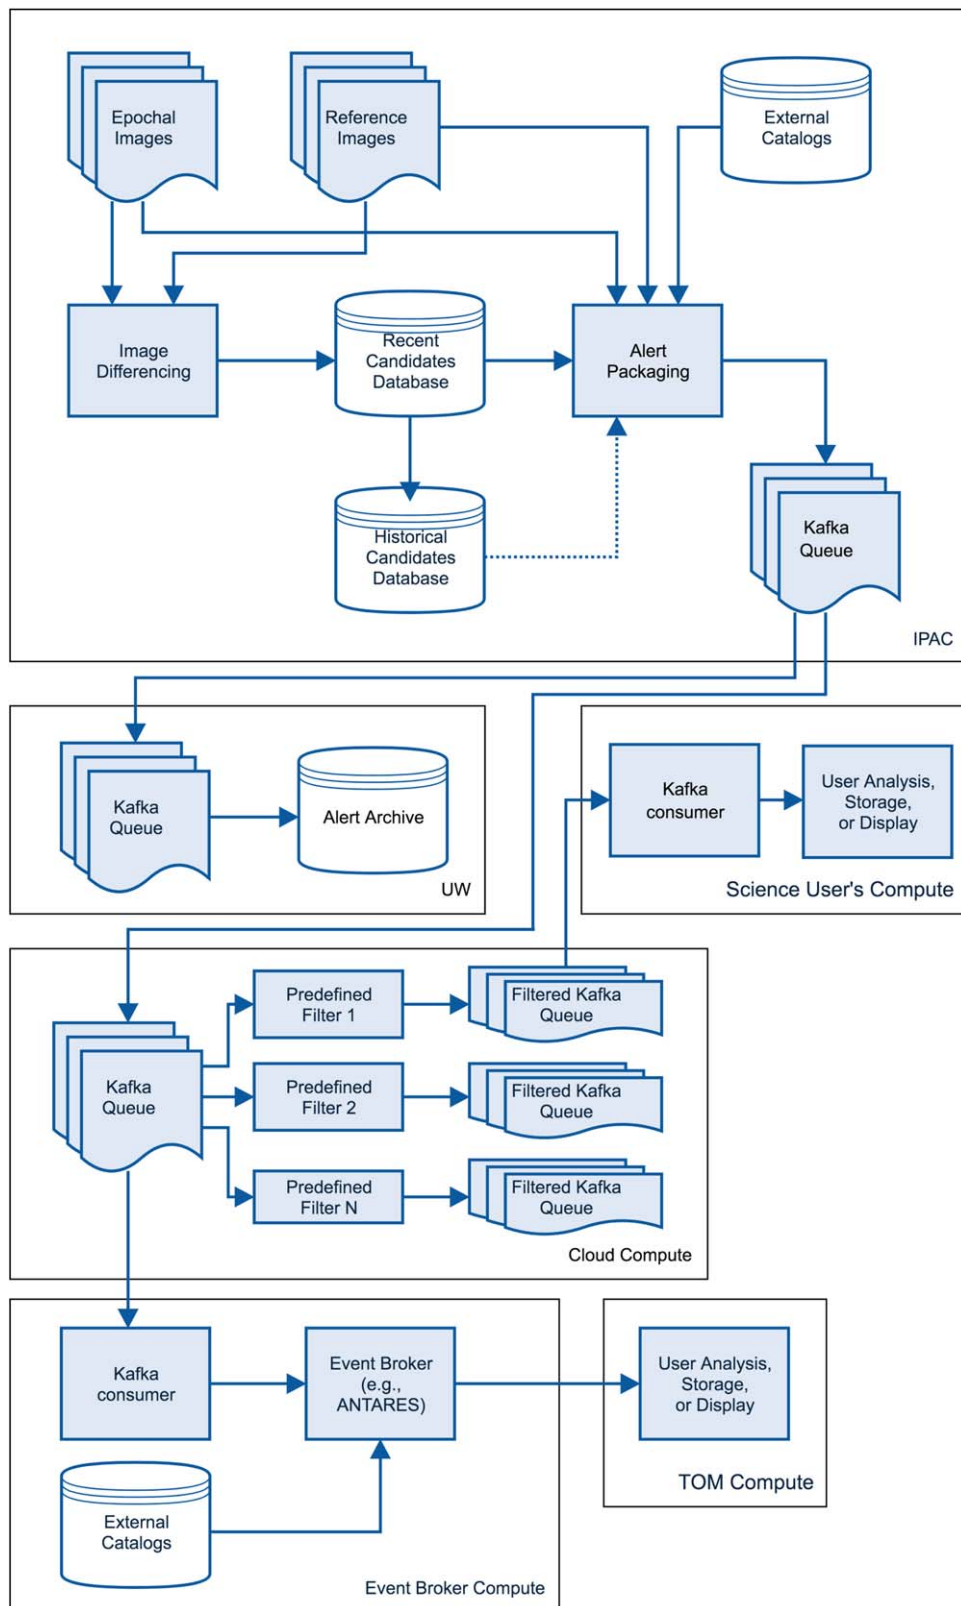
ZADS is in the process of deploying pre-defined filters of the stream for a number of science use cases. These pre-defined filters are to be deployed using Python code in Docker containers. Each filter reads from a Kafka topic and produces to a separate Kafka topic. Alerts are pulled into a filter and deserialized from Avro into Python dictionaries. The dictionaries then pass through the if-then logic of a simple Python function that returns True if the alert is of interest and False if not. Alerts that pass the filter are then serialized again to Avro and pushed back to a Kafka topic, for downstream consumption.

The above process can also be employed for user-defined filters. In order to build reasonable working filters, a user might find a sample set of batch alerts useful. A batch of alerts can be collected by consuming from the stream and writing each Avro packet out to files on disk. The batched Avro packets can be processed with the same filtering code used to process alert Avro packets received in the live stream, given that the data structure is the same for both batch and stream processing. Users can then develop filtering code using batch data as a training set and deploy filters on the live stream with few changes to the code.

## 4. ZADS Deployment in ZTF Nightly Processing

A diagram of the flow of ZTF alert data through the ZADS pipeline is shown in Figure 2. The nightly alert pipeline begins at IPAC at Caltech when new detections are produced from image differencing between epochal and reference image data. The candidates are loaded into a database where they are cross-matched to a database of previous historical ZTF candidates

<sup>12</sup> <https://www.docker.com>



**Figure 2.** Diagram of the flow of alert data through the ZTF Alert Distribution System. (A color version of this figure is available in the online journal.)



and to the Pan-STARRS1 and *Gaia* catalogs. Alerts are packaged into Avro format, containing, e.g., brightness, time, position, image cutouts of before, event, and subtraction, cross-matching to other catalogs, and previous detection information. The full schema as well as sample alert packets are available.<sup>13</sup> Table 1 provides a high-level summary of some of the fields included in the alert packets. The pipeline produces Avro alerts in parallel, using Kafka as a buffer, reducing to fewer streams, or Kafka topics, divided by date and by program ID. This separates internal project collaboration alerts and public alerts into separate streams, which are then ready for further distribution outside of IPAC.

Alerts from moving objects are included in the stream, but Minor Planet Center reporting is handled separately from ZADS, as ZADS is a general purpose alert distribution mechanism and interfaces with science-specific interfaces require additional filtering and post-processing. Candidate moving objects are submitted to the MPC after a Moving Object Processing System process runs. More detail on moving object handling is provided in the ZTF Data System overview paper (Masci et al. 2018).

From IPAC, alerts are pulled into another instance of Kafka at the University of Washington (UW) via Kafka’s Mirror-Maker functionality. The UW Kafka instance runs as a small cluster of three Kafka brokers in Docker containers, mirroring collaboration and public program alert stream topics. From the UW Kafka instance, a local consumer reads each night of data and saves the stream to an alert archive. This archive is primarily for backup purposes and to provide limited access to past alerts in case of connection issues with community brokers. While more capable community broker tools are in development, we are also using this archive to provide community access to public alerts as bulk downloads.<sup>14</sup> The UW Kafka instance also distributes alerts to ZTF internal collaboration partners.

The IPAC cluster also feeds the public stream into a cloud-based Kafka instance, populated again by using the Mirror-Maker functionality. This Kafka instance is currently running on the DigitalOcean cloud infrastructure. The public stream is then available to astronomical event “community brokers,” Target and Observation Managers (TOMs), and similar systems through the cloud-based Kafka instance. Both the cloud-based Kafka broker and the UW Kafka broker keep a 7 day “hot cache” of alert streams, making one week of past alerts available to community brokers in addition to the live stream. Alert stream topics are separated into 16 partitions. This allows consumers of the stream to parallelize readers into a group of up to 16, allowing faster reading and parallel downstream filtering and processing.

The cloud-based Kafka system will also feed into pre-defined community filters. These pre-defined community filters will run co-located in the cloud. Community filters are intended to be a first pass at detecting and classifying, for example, supernova or variable stars. Alerts that pass a community filter are routed back to the Kafka instance in a separate topic per filter. These filtered streams are then available for analysis by science users.

## 5. Performance and Metrics

ZTF produces around 600,000 to 1.2 million alerts on a full night of observing. The ZADS system has been observed to withstand a throughput of over 2 million alerts with no technical issues. The size of each individual alert is about 60 KB, dominated by the size of the included cutout images. The total volume of ZTF alerts pushed through ZADS over one night can then amount to over 70 GB of data.

The total end-to-end ZADS pipeline results in alerts available roughly 20 minutes after a field is observed, dominated by the data reduction time. The median time between alert candidate production and packaging to availability in the IPAC Kafka instance is about 6 s. In simulations, the serialization of 1000 ZTF alerts into Avro format and submission to a Kafka topic, and transfer to the consumer was observed to be 4.2 s. Data transfer rates of about 100 MBps to the cloud have been observed. This equates to over 80k alerts/minute or 6.6k alerts in 5 s.

## 6. Interface

ZADS provides two main endpoints for scientific users to connect and receive alerts. One way to access the public stream is via astronomical event “community brokers.” Community brokers consume the full public stream and provide value added services such as cross-matching with external catalogs to enrich alerts. After community broker processing and additional filtering, scientific users can access these enriched alerts via the broker interfaces.

Science users may also wish to directly access the public stream from their own compute resources. At present, access to the public stream is expected to be limited to filtered topics from the pre-defined community filters. Example code demonstrating how an approved user can connect directly to a pre-defined community filtered stream as a Kafka topic is available at [https://github.com/ZwickyTransientFacility/alert\\_stream](https://github.com/ZwickyTransientFacility/alert_stream). This code could also be used with any ZTF alert distribution hub that utilizes a Kafka instance to stream alerts. Note that this code is only an example and will not allow unknown users to connect to ZADS. Connecting directly to ZADS is only available to broker projects who have signed MOU agreements. Access to the stream is currently restricted by IP address. Encryption, authentication, and additional restrictions on consumers (topic access controls and throttling

<sup>13</sup> <https://github.com/ZwickyTransientFacility/ztf-avro-alert>

<sup>14</sup> <https://ztf.uw.edu/alerts/public/>

stream consumption) are available in Kafka and planned to be implemented. These features will be essential in future projects particularly with many consumers.

## 7. Summary

We have described here the ZTF Alert Distribution System and the technologies used to implement the pipeline. ZTF Avro formatted alerts are streamed from IPAC to the University of Washington as well as a cloud computing service and subsequently to downstream community brokers and scientific users via the Apache Kafka data streaming platform. Along the pipeline, alerts are collected to an alert archive and also processed with simple pre-defined community filters as a first pass at creating filtered streams separated into various distinct astronomical sources.

The ZADS platform has successfully transmitted more alerts than have ever been distributed via VOEvents, on the order of one million alerts nightly. This is the first successful demonstration of at-scale alert distribution in an astronomical context, using industry standard tools and serialization formats, and may provide guidance for the evolution of astronomical standards in preparation for the Large Synoptic Survey Telescope. The alert volume of ZTF approaches 10% of LSST, with significantly more modest hardware used for ZADS; we therefore feel this is strong evidence to the LSST alert problem being tractable. The robustness of the distribution system and filtering capabilities make ZADS promising as a precursor to the high volume of alerts expected from the upcoming LSST.

Based on observations obtained with the Samuel Oschin Telescope 48 inch and the 60 inch Telescope at the Palomar Observatory as part of the Zwicky Transient Facility project, a scientific collaboration among the California Institute of Technology, the Oskar Klein Centre, the Weizmann Institute of Science, the University of Maryland, the University of Washington, Deutsches Elektronen-Synchrotron, the University of Wisconsin-Milwaukee, and the TANGO Program of the University System of Taiwan. Further support is provided by the U.S. National Science Foundation under grant No. AST-1440341.

M. Patterson, E. Bellm., and K. S. Krughoff acknowledge support from the Large Synoptic Survey Telescope program, which is supported in part by the National Science Foundation through Cooperative Agreement 1258333 managed by the Association of Universities for Research in Astronomy (AURA), and the Department of Energy under contract No.











DE-AC02-76SF00515 with the SLAC National Accelerator Laboratory. Additional LSST funding comes from private donations, grants to universities, and in-kind support from LSSTC Institutional Members.

M. Patterson, E. Bellm, V. Z. Golkhou, and M. Juric acknowledge support from the University of Washington College of Arts and Sciences, Department of Astronomy, and the DIRAC Institute. University of Washington's DIRAC Institute is supported through generous gifts from the Charles and Lisa Simonyi Fund for Arts and Sciences, and the Washington Research Foundation.

E. Bellm is supported in part by the NSF AAG grant 1812779 and grant #2018-0908 from the Heising-Simons Foundation.

M. Juric acknowledges the support of the Washington Research Foundation Data Science Term Chair fund, and the UW Provost's Initiative in Data-Intensive Discovery.

## ORCID iDs

Maria T. Patterson  <https://orcid.org/0000-0002-4753-3387>  
 Eric C. Bellm  <https://orcid.org/0000-0001-8018-5348>  
 Ben Rusholme  <https://orcid.org/0000-0001-7648-4142>  
 Frank J. Masci  <https://orcid.org/0000-0002-8532-9395>  
 Mario Juric  <https://orcid.org/0000-0003-1996-9252>  
 K. Simon Krughoff  <https://orcid.org/0000-0002-4410-7868>  
 V. Zach Golkhou  <https://orcid.org/0000-0001-8205-2506>  
 Matthew J. Graham  <https://orcid.org/0000-0002-3168-0139>  
 Shrinivas R. Kulkarni  <https://orcid.org/0000-0001-5390-8563>  
 George Helou  <https://orcid.org/0000-0003-3367-3415>

## References

- Allan, A., Denny, R. B., & Swinbank, J. D. 2017, arXiv:1709.01264  
 Bellm, E., Kulkarni, S., Graham, M. J., et al. 2018, PASP, in press  
 Cao, Y., Kulkarni, S. R., Howell, D. A., et al. 2015, *Natur*, **521**, 328  
 Cenko, S. B., Urban, A. L., Perley, D. A., et al. 2015, *ApJL*, **803**, L24  
 Gal-Yam, A., Arcavi, I., Ofek, E. O., et al. 2014, *Natur*, **509**, 471  
 Graham, M. J., Kulkarni, S., Bellm, E., et al. 2018, PASP, submitted  
 Ivezić, Z., Tyson, J. A., Abel, B., et al. 2008, arXiv:0805.2366  
 Kreps, J., Narkhede, N., Rao, J., et al. 2011, in Proc. NetDB (Athens), 1  
 Maeda, K. 2012, in 2012 Second Int. Conf. Digital Information and Communication Technology and its Applications (DICTAP), 177  
 Masci, F., Laher, R., Bellm, E., et al. 2018, PASP, in press  
 Narayan, G., Zaidi, T., Soraisam, M. D., et al. 2018, *ApJS*, **236**, 9  
 Rutledge, R. E. 1998, *PASP*, **110**, 754  
 Seaman, R., Williams, R., Allan, A., et al. 2011, <http://www.ivoa.net/documents/VOEvent/20110711/>  
 Staley, T. D., & Fender, R. 2016a, <https://4pisky.org/voevents/>  
 Staley, T. D., & Fender, R. 2016b, arXiv:1606.03735  
 Williams, R. D., & Seaman, R. 2006, in ASP Conf. Ser. 351, Astronomical Data Analysis Software and Systems XV, ed. C. Gabriel et al. (San Francisco, CA: ASP), 637